

Project - High Level Design

On

Multi-Agent Manufacturing System

Course Name: Agentic AI

Institution Name: Medicaps University – Datagami Skill Based Course

Student Name(s) & Enrolment Number(s):

Sr no	Student Name	Enrolment Number
01	Om Tiwari	EN22CS301669
02	Paridhi Shirwalkar	EN22CS301684
03	Nitesh Chourasiya	EN22CS301660
04	Mradul Jain	EN22CS301616

Group Name: Group 01D6

Project Number: AAI-12

Industry Mentor Name:

University Mentor Name: Nishant Shrivastav

Academic Year: 2025-2026

Table of Contents

1. Introduction.
 - 1.1. Scope of the document.
 - 1.2. Intended Audience
 - 1.3. System overview.
2. System Design.
 - 2.1. Application Design
 - 2.2. Process Flow.
 - 2.3. Information Flow.
 - 2.4. Components Design
 - 2.5. Key Design Considerations
 - 2.6. API Catalogue.
3. Data Design.
 - 3.1. Data Model
 - 3.2. Data Access Mechanism
 - 3.3. Data Retention Policies
 - 3.4. Data Migration
4. Interfaces
5. State and Session Management
6. Caching
7. Non-Functional Requirements
 - 7.1. Security Aspects
 - 7.2. Performance Aspects
8. References

1. Introduction

1.1 Scope of the Document

This document defines the low-level design for a Multi-Agent Manufacturing System that automates supplier sourcing and generates a structured supplier comparison report. The system uses a collaborative, role-based agent architecture with a strict handoff protocol to ensure predictable outputs.

The scope includes:

- System workflow and execution lifecycle
- Component-level architecture and responsibilities
- Data design and schemas for inter-agent handoffs
- Interfaces and APIs for UI and internal modules
- State/session handling, caching, and persistence
- Non-functional requirements (performance, security, reliability)
- Error handling, retry logic, and validation strategy

Out of scope:

- Production-grade deployment and autoscaling setup
- Paid integrations like enterprise vendor databases
- Advanced RAG pipelines (optional future enhancement)

1.2 Intended Audience

This document is intended for:

- Developers implementing the system (Python + CrewAI)
- Reviewers evaluating the architecture (LLD review)
- Testers validating data integrity and workflow states
- Maintainers extending agents/tools or adding new workflows

1.3 System Overview

The system is composed of two specialized agents orchestrated through a central controller:

- **Researcher Agent:** collects supplier information based on a user query and produces a raw dataset.
- **Writer Agent:** transforms the raw dataset into clean, structured supplier records and generates a formatted comparison report.

- **Orchestrator:** manages the end-to-end flow, validation, retries, state persistence, and artifact creation.
- **UI Layer (Streamlit):** provides input capture, model/API key handling, execution trigger, and output visualization/download.

Key design principles:

- Strict JSON-based handoff artifacts between agents
- Schema validation at each stage to prevent garbage-in → garbage-out
- Session-based runs with deterministic storage paths
- Defensive orchestration with controlled retries and fallbacks

2. System Design


2.1 Application Design


Architecture style: Layered + Orchestrated multi-agent pipeline.

Layers


1. Presentation Layer (Streamlit UI)
 - Input form (manufacturing query)
 - Model selection + API key handling
 - Progress logging and output rendering
2. Application Layer (Orchestrator + Workflow Engine)
 - Creates a unique run/session
 - Executes Research crew then Writer crew
 - Validates outputs, stores artifacts, marks workflow state
3. Domain Layer (Agent Logic + Tasks)
 - Agent definitions and role prompts
 - Task templates and output format constraints
4. Data Layer (Storage + Schemas)
 - JSON schemas for handoff
 - File-based persistence and caching

Multi-Agent Workflow Activity





Backend Class Structure



2.2 Process Flow

Primary Flow (Happy Path)

1. User submits supplier sourcing query via UI
2. Orchestrator creates run_id and initializes workflow state
3. Researcher Agent runs and returns Raw Supplier Dataset (JSON)
4. Orchestrator validates raw JSON (schema + minimal quality checks)
5. Writer Agent consumes raw dataset and returns:
 - o Structured Supplier Records (JSON)
 - o Supplier Comparison Report (Markdown/HTML/Text)
6. Orchestrator stores outputs, finalizes state, returns result to UI
7. UI renders report and provides downloadable artifacts

Failure Handling (High-Level)

- If Researcher fails → retry with backoff (configurable)
- If schema invalid → request agent correction (limited attempts) or fail run with detailed error artifact
- If Writer fails → retry; if still fails, store raw research artifact for debugging and return partial output

2.3 Information Flow

Input → Output data movement

- UI sends: {query, api_key, model}
- Orchestrator sends to research crew: {query, run_id}
- Researcher returns: raw_suppliers.json (untrusted)
- Orchestrator validates + stores: raw_suppliers_validated.json
- Writer consumes validated raw: generates:
 - o structured_suppliers.json
 - o report.md or report.txt
- Orchestrator returns a summary object to UI:
 - o run_id
 - o artifact paths
 - o report preview snippet

- status + timestamps

2.4 Components Design

2.4.1 Streamlit Frontend

Responsibilities:

- Collect query text (supplier sourcing request)
- Allow selection of model (Gemini, etc.)
- Execute orchestrator run
- Render results: report, tables, JSON preview
- Provide download buttons for artifacts

Key UI states:

- Idle
- Running
- Completed
- Failed (show error + debug artifact link)

2.4.2 Orchestrator (Core Controller)

Responsibilities:

- Create run/session ID
- Prepare storage directories
- Execute agent crews sequentially
- Validate and normalize outputs
- Save artifacts and workflow metadata
- Return structured results to UI

Important orchestration decisions:

- Sequential is preferred for deterministic handoff correctness
- Strict schema validation occurs between agents
- Adds guardrails: max retries, max tokens, timeouts (where supported)

2.4.3 Researcher Agent

Responsibilities:

- Understand the sourcing request
- Identify supplier candidates

- Collect key supplier attributes (name, location, products, certifications, lead time, MOQ, pricing ranges if possible, contact channels)
- Produce a raw dataset that is complete enough for synthesis

Output format:

- JSON list of supplier objects
- Includes source references (links or notes)
- Includes confidence/quality notes per supplier if applicable

2.4.4 Writer Agent

Responsibilities:

- Clean and standardize supplier data
- Fill missing fields with "unknown" or null rather than hallucinating
- Rank suppliers based on scoring logic (configurable)
- Produce:
 - structured normalized JSON
 - human-readable comparison report with findings, risks, and gaps

Report sections:

- Executive summary
- Comparison table
- Ranking explanation
- Risks/gaps (missing certifications, unclear lead times, weak contact info)
- Next steps and recommended outreach questions

2.4.5 Storage Layer

Responsibilities:

- Write artifacts to disk under run-specific folder
- Maintain workflow state file
- Optional caching for identical queries

Folder example:

```
artifacts/  
  run_20260219_101530/  
    input.json  
    state.json
```

raw_suppliers.json
raw_suppliers_validated.json
structured_suppliers.json
report.md
logs.jsonl

2.5 Key Design Considerations

2.5.1 Deterministic Outputs

- Force agents to return JSON in strict schema
- Validate using Pydantic/JSONSchema
- Normalize fields (string trimming, list coercion, default values)

2.5.2 Preventing Hallucinated Data

- Writer must not fabricate missing supplier details
- Use explicit rules:
 - If unknown: set "unknown"
 - If not found: set null
 - Add missing_fields list per supplier

2.5.3 Observability

- Store run logs with timestamps
- Capture agent prompts and responses (optional, sanitized)
- Maintain state.json transitions for audit

2.5.4 Extensibility

- Add new agents later (e.g., Negotiator Agent, Compliance Agent)
- Add tools (web search, vector DB, document ingestion)
- Add new workflows (RFQ generation, supplier onboarding checklist)

2.6 API Catalogue

2.6.1 Frontend → Orchestrator (Internal Call)

Function: run(query: str) -> dict

Request:

{

```
"query": "Find aluminum die casting suppliers in India for automotive parts"
}

Response:

{
  "run_id": "run_20260219_101530",
  "status": "completed",
  "artifacts": {
    "raw": "artifacts/run_.../raw_suppliers.json",
    "structured": "artifacts/run_.../structured_suppliers.json",
    "report": "artifacts/run_.../report.md"
  },
  "summary": "Top suppliers ranked by certifications + lead time..."
}
```

2.6.2 Storage API (Internal)

- init_run() -> run_id
- save_json(run_id, filename, data)
- load_json(run_id, filename)
- save_text(run_id, filename, text)
- update_state(run_id, patch)

3. Data Design

3.1 Data Model

3.1.1 Raw Supplier (Researcher Output)

Core fields:

- supplier_name (string)
- country / city (string)
- capabilities (list of strings)
- industries_served (list)
- certifications (list)
- lead_time (string)

- moq (string/number)
- pricing_info (string, optional)
- contact (object: email/phone/website)
- sources (list: url or notes)

3.1.2 Structured Supplier (Writer Output)

Adds:

- normalized fields (standard formats)
- scoring fields (quality_score, cost_score, lead_time_score)
- rank (integer)
- gaps (list)
- confidence (low/medium/high)

3.2 Data Access Mechanism

- File-based storage (JSON + markdown) for simplicity and portability
- Each run is isolated by run_id
- Reads happen only within the run directory
- Optional caching layer:
 - Hash of query + model → stored response reuse

3.3 Data Retention Policies

- Default retention: keep artifacts locally for debugging/demo
- Optional cleanup job:
 - Delete runs older than N days
 - Keep only report.md + structured_suppliers.json

Sensitive data considerations:

- API keys should never be written to disk
- If storing prompts/responses, strip secrets and tokens

3.4 Data Migration

If schema changes:

- version your schema: schema_version field in each JSON artifact
- maintain migration scripts:
 - v1 → v2: rename fields, add defaults, convert types

- backward compatibility rules:
 - Writer can accept v1 raw schema and output v2 structured schema

4. Interfaces

4.1 User Interface

- Text area: sourcing query
- Inputs: API key, model name
- Buttons: Run, Reset, Download report/JSON
- Output panels:
 - Status (Running/Done/Failed)
 - Report preview
 - Supplier table

4.2 Internal Module Interfaces

- Agents module exposes: build_researcher(), build_writer()
- Tasks module exposes: research_task(agent), write_task(agent)
- Schemas module exposes: Pydantic models + validators
- Orchestrator exposes: run(query)

5. State and Session Management

5.1 Run Lifecycle

State machine example:

- CREATED
- RESEARCH_RUNNING
- RESEARCH_DONE
- WRITE_RUNNING
- COMPLETED
- FAILED

State file (state.json) fields:

- run_id
- status

- timestamps (created_at, updated_at)
- error (if failed)
- artifacts list

5.2 Session Isolation

- Each run gets a unique directory
- No shared mutable state between runs
- Thread safety: avoid global variables for run data

6. Caching

6.1 Caching Strategy

Goal: reduce cost and speed up repeat queries.

Cache key:

- hash(query + model + prompt_version)

Cache store:

- cache/index.json maps key → artifact path

Cache rules:

- Only cache if run completed successfully
- Invalidate cache if schema_version changes or prompt_version changes

7. Non-Functional Requirements

7.1 Security Aspects

- Never store API keys on disk
- .env included in .gitignore
- Validate and sanitize user input (length limits, safe characters)
- Avoid executing user input as code
- If adding web tools:
 - restrict allowed domains (optional)
 - log sources used

7.2 Performance Aspects

- Target: usable demo experience on a laptop

- Minimize LLM calls:
 - one research run + one writer run (baseline)
 - limit agent iterations
- Implement backoff retries for transient errors (429, network)
- Graceful degradation:
 - return partial results if writer fails but research succeeded

Reliability targets:

- Each run produces a final state.json even if failed
- Errors stored as error.json for debugging

8. References

- CrewAI documentation (agents, tasks, crews, orchestration concepts)
- LLM provider API documentation (Gemini/OpenAI depending on provider)
- Pydantic / JSON Schema validation references
- Streamlit documentation for UI patterns and session handling